

Robot Learning: A Tutorial

Francesco Capuano  ... Adil Zouitine  Pepijn Kooijmans  Thomas Wolf  Michel Aractingi 

 École Normale Supérieure Paris-Saclay,  Hugging Face

Abstract



Hugging Face

Recent advancements in foundation models have enabled unprecedented progress in robotics, particularly in the development of increasingly more capable generalist robotics models. Leveraging modern learning-based approaches, the robotics community produced the first general models for tasks including locomotion and manipulation, and many said results have been openly released to the open-source community. The premise of generalist models is to learn control strategies across tasks and different robot embodiments, a long-standing challenge in robot learning. This work aims at being a tutorial on the most common techniques used for modern robot learning, covering both the conceptual underpinnings of the most prominent approaches, and pairing them with reproducible implementations using `lerobot` (<https://github.com/huggingface/lerobot>), the full-stack open-source robotics library developed by Hugging Face. This tutorial is designed to serve as a self-contained, structured reference for robotics practitioners working with `lerobot`, as well as a practical guide for researchers in the field of robot learning.

1 Introduction

1.1 Motivation: The Interdisciplinary Nature of Robotics in the Machine Learning Era

Robotics in 2025 sits at the intersection of classical model-based control, machine perception, and large-scale machine learning. Foundational problems in locomotion, manipulation, and whole-body control demand reasoning across rigid-body dynamics, contact modeling, planning under uncertainty, and high-dimensional function approximation. At the same time, end-to-end learning has matured from proof-of-concept demonstrations to systems that benefit from internet-scale multimodal pretraining and robotics-specific fine-tuning, closing the gap between laboratory benchmarks and deployment in unstructured settings. This tutorial embraces that interdisciplinarity: it treats modern robot learning as a synthesis of classical priors and data-driven policies rather than a replacement of one by the other.

1.2 Scope and Contributions of This Work

This document serves two purposes. First, it provides a concise, academically rigorous overview of core concepts in classical robotics that remain essential for principled system design (kinematics, dynamics, planning, and control). Second, it surveys contemporary learning-based methods for robotic control, with an emphasis on (i) reinforcement learning (RL) and imitation learning (IL), (ii) single-task policy architectures such as transformer-based action chunking and diffusion policies, and (iii) multi-task vision–language–action (VLA) models. Throughout, we complement exposition with *reproducible* implementation guidance using the open-source `LeRobot` framework, to lower the barrier between concept and practice while maintaining scientific rigor.

1.3 Structure of the Report

Section 2 reviews classical robotics foundations and their limitations in contact-rich, high-DOF regimes. Section 3 formulates learning-based control via RL and IL, highlighting problem setups, algorithms, and known failure modes (e.g., reward misspecification, simulation gaps, and safety constraints). Section 4 details single-task policy families (transformer chunking, diffusion) and their practical training recipes. Section 5 surveys multi-task VLA models (e.g., RT-1/RT-2, OpenVLA, $\pi_0/\pi_{0.5}$, and SmolVLA), together with integration patterns and experimental evaluation protocols. Section 6, not covered here, discusses emerging directions (e.g., world models and post-training) beyond the present scope.

2 Generating Motion

Robotics is concerned with producing real-world motion in a way that is useful, reliable and safe. At its core, robotics is thus an inherently multidisciplinary domain, as producing motion requires interfacing various hardware and software components, each at various levels. As a direct consequence of its multi-disciplinary nature, robotics traditionally employed an array of methods concerned with its main purpose of producing motion. Methods to produce robotics motion range from traditional *explicit* models—leveraging precise descriptions of the mechanics of robots’ rigid bodies and their interactions with eventual obstacles in the robots environment—to fully *learning-based* systems, offloading modelling the robot mechanics and rather treating motion as a statistical pattern to learn given multiple sensorimotor readings. Between these two extrema, a variety of methods exist, with many borrowing traits from the opposite extremum to tackle specific challenges. On the one hand, as learning-based systems can benefit from information relative to the physics of particularly static tasks, Temporal Difference (TD) methods have been complemented with Model-Predictive Control (MPC) (?). On the other hand, explicit models may be relying on assumptions proving overly simplistic—or even unrealistic—in practice, as in the case of assuming perfect observability of a robot state using sensorimotor readings. In this context, feedback loops at the control level help mitigate the effects of poor state estimation, complementing the motion planning process with discrepancy-from-target information, similarly to how loss functions are used in statistical learning. Figure ?? graphically illustrates the most prominent models in general-purpose applications, and we refer the interested reader to ? for a comprehensive overview of both general and application-specific methods for motion generation. Aiming at introducing the inherent benefit of learning-based approaches in the context of an increasing availability of robot-data—the main focus of this tutorial—in the rest of this section we will present an illustrative example for the task of manipulation using a rather simplistic technique in the broader context of traditional robotics: inverse kinematics. While far from more advanced techniques that have proven groundbreaking in specific real-world applications (among many others, (??)), we believe the basics primitive of a full description of movement provide sufficient intuition as to the motivation of learning-based methods in the context of modern robotics.

2.1 A Toy Example: (Planar) Manipulation

At its core, robotics deals with producing motion via actuating joints connecting nearly entirely-rigid links. Based on whether the generated motion modifies (1) the relative state of the robot with respect to its environment, or (2) the absolute state of the environment.

Toy example of planar robot

Full physical description by means of forward kinematics to generate movement

Most of the times you want to track something, so you do inverse kinematics

What is inverse kinematics

2.1.1 Overcoming Estimation Error via Feedback Loops

Adding feedback loops to a (simplified) dynamical system

Tuning the gain is not immediate and rather cumbersome

2.2 Limitations

Reinforcement Learning for pivoting task’s limitations for robot learning

3 Learning-Based Approaches to Robotics

3.1 Reinforcement Learning (RL) for Robotic Control

3.1.1 Problem Formulation and Control Objectives

We pose control as a Markov decision process (MDP) with state space \mathcal{S} , action space \mathcal{A} , transition kernel P , reward r , and discount γ . A policy $\pi_\theta(a | s)$ maximizes $J(\theta) = \mathbb{E}[\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t)]$. Practical objectives include setpoint tracking, task success, energy regularization, and safety constraints (e.g., joint limits, collision avoidance).

3.1.2 Policy Optimization Methods in Robotics

Modern practice spans on-policy (e.g., TRPO/PPO) and off-policy (e.g., DDPG/TD3/SAC) actor-critic methods, as well as model-based RL that learns or exploits dynamics for sample efficiency. Off-policy approaches are favored for real robots due to data efficiency and the ability to incorporate replay. Auxiliary techniques—domain randomization, dynamics identification, and constraint handling—facilitate sim-to-real transfer and safe exploration.

3.1.3 Practical Implementation: Training RL Policies with LeRobot

LeRobot provides standardized dataset and model abstractions, and can be integrated with established RL backends. A typical workflow comprises: (i) environment and observation/action specification (simulated or real), (ii) policy class selection (e.g., SAC/PPO wrappers), (iii) logging and evaluation hooks, (iv) sim-to-real calibration (actuation limits, latency compensation), and (v) deployment with safety interlocks. The framework’s dataset utilities simplify offline RL pretraining or behavior cloning initialization before online fine-tuning.

3.1.4 Limitations of RL in Real-World Robotics: Simulators and Reward Design

Notwithstanding successes in quadruped locomotion and manipulation, practical bottlenecks persist: expensive data collection, reward misspecification and shaping burden, partial observability, and safety during exploration. High-fidelity simulation reduces real-world trials but introduces transfer gaps; reward engineering often requires domain expertise and iterative tuning. Stable, sample-efficient, and safety-aware RL remains an active research frontier.

3.2 Imitation Learning (IL) for Robotics

3.2.1 Leveraging Real-World Demonstrations

IL bypasses online reward optimization by learning from expert trajectories. In visuomotor settings, demonstrations pair high-dimensional observations (RGB, depth, proprioception) with action sequences. Policy learning reduces to supervised learning under covariate shift, necessitating strategies that mitigate compounding errors.

3.2.2 Reward-Free Training and Data-Centric Perspectives

Data-centric IL emphasizes broad, diverse, and real-world demonstration corpora to improve generalization. Pretraining on heterogeneous demonstrations and fine-tuning on task-specific data can produce robust policies without explicit rewards, especially for contact-rich manipulation.

3.2.3 A Taxonomy of IL Approaches

Behavior Cloning (BC) minimizes one-step imitation loss; interactive IL (e.g., DAgger) aggregates data under the learner’s state distribution; adversarial/apprenticeship methods recover reward surrogates; and diffusion- or flow-based action generators model multimodal action distributions. Interactive feedback and on-the-fly corrections further reduce covariate shift in deployment.

4 Single-Task Policy Architectures

4.1 Action Chunking with Transformers

4.1.1 Model Architecture and Training Objectives

Action Chunking with Transformers (ACT) models short horizons of low-level actions as fixed-length “chunks” conditioned on recent observations and proprioception. A transformer parameterizes a distribution over action sequences, enabling parallel prediction of temporally coherent control signals. Supervised training on teleoperated demonstrations minimizes sequence-level losses (e.g., mean-squared error on joint velocities/torques), optionally with scheduled sampling or consistency regularization to improve rollout stability.

4.1.2 Practical Implementation in LeRobot

A practical ACT pipeline comprises (i) synchronized demonstration collection (vision and robot states) with precise timestamping, (ii) dataset serialization into a standardized episodic format, (iii) chunked windowing of action targets,

(iv) transformer training with early stopping on validation rollouts, and (v) deployment with action-rate limiting and safety monitors. Careful calibration of control frequency and chunk length is critical for fine manipulation.

4.2 Diffusion-Based Policy Models

4.2.1 Generative Modeling for Action Sequences

Diffusion policies treat action generation as conditional denoising: starting from noise in action space, a learned score model iteratively refines to produce feasible control sequences conditioned on observations (and optionally goals or language). This formulation naturally captures multimodality (multiple valid strategies), scales to higher-dimensional actions, and supports receding-horizon execution.

4.2.2 Practical Implementation

Training proceeds by corrupting ground-truth action sequences and fitting a time-indexed network to denoise under teacher forcing; inference uses a small number of reverse steps with horizon MPC-style rollouts. In practice, datasets benefit from action normalization, viewpoint augmentation, and careful tuning of denoising steps versus control latency. Within LeRobot, diffusion policy trainers can be paired with standardized dataset loaders and evaluation callbacks to track success rates and robustness under perturbations.

5 Multi-Task Policies: Vision–Language–Action (VLA) Models in Robotics

5.1 Overview of Major Architectures: RT-1, RT-2, OpenVLA, π_0 , $\pi_{0.5}$, SmolVLA

VLA models integrate perception, instruction following, and action generation within a unified network. RT-1 demonstrated large-scale, robot-collected datasets and transformer policies capable of executing hundreds of distinct manipulation tasks; RT-2 extended this idea by co-training with web-scale vision–language corpora and representing actions as discrete tokens, improving semantic generalization to novel objects and instructions. OpenVLA adopts an open 7B backbone that fuses pretrained vision encoders with a language model and trains on large, diverse robot demonstrations, reporting strong cross-embodiment performance and efficient fine-tuning. The π_0 family introduces flow-matching action heads atop a pretrained VLM to produce continuous, high-frequency motor commands for dexterous skills; $\pi_{0.5}$ augments co-training for broader open-world generalization. Complementing these, SmolVLA targets efficiency and accessibility with a compact ($\sim 4.5 \times 10^8$ parameters) architecture and an interleaved attention “action expert” for low-latency control.

5.3 Practical Implementation: Integrating VLAs with LeRobot

Integration typically follows three phases. *Pretraining*: initialize the VLM backbone (e.g., SigLIP/DINOv2 features) and attach an action head (autoregressive, diffusion, or flow-matching) trained on heterogeneous demonstrations curated in a standardized format. *Task adaptation*: fine-tune with instruction-augmented, robot-specific data using low-rank adaptation where possible; calibrate action scaling, control-rate chunking, and observation encodings. *Deployment*: quantize or distill for on-robot inference; implement asynchronous inference stacks to decouple perception and control loops; add safety fallbacks and confidence gating. LeRobot dataset and model utilities simplify all three phases by providing consistent IO, evaluation tools, and model hosting.

5.4 Experimental Evaluation

A principled evaluation protocol should measure (i) success rate across held-out tasks and objects, (ii) language grounding fidelity under paraphrases and compositional instructions, (iii) cross-embodiment transfer (train/test robot mismatch), and (iv) robustness to distribution shift (lighting, clutter, distractors). Where possible, report zero-shot and few-shot adaptation performance, ablate pretraining datasets and backbones, and include real-world trials with fixed seeds and video evidence. Public, diverse datasets and standardized success criteria are critical to comparable results.

References